

Scale search powered apps with Elasticsearch, k8s and go

Maxime Boisvert



2016

Was a major year for me

Joined Shopify

Had a son

2017

Joined the search infrastructure team



You searched for: ***cap***

We found 27 results


Search again

cap



Applications moving from various cloud
provider and data centers to Google
Cloud





Growing number of Elasticsearch
clusters



Search operator

Based on Elasticsearch, Kubernetes,
Golang and Docker

Elasticsearch basics

- Cluster
- Node
- Index
- Replicas

Kubernetes basics

- Pod
- Deployment
- Statefulset
- Service
- Configmap
- Custom Resource Definition

Good practices are not always in place



Simple configuration

Simple configuration: elasticsearch.yml

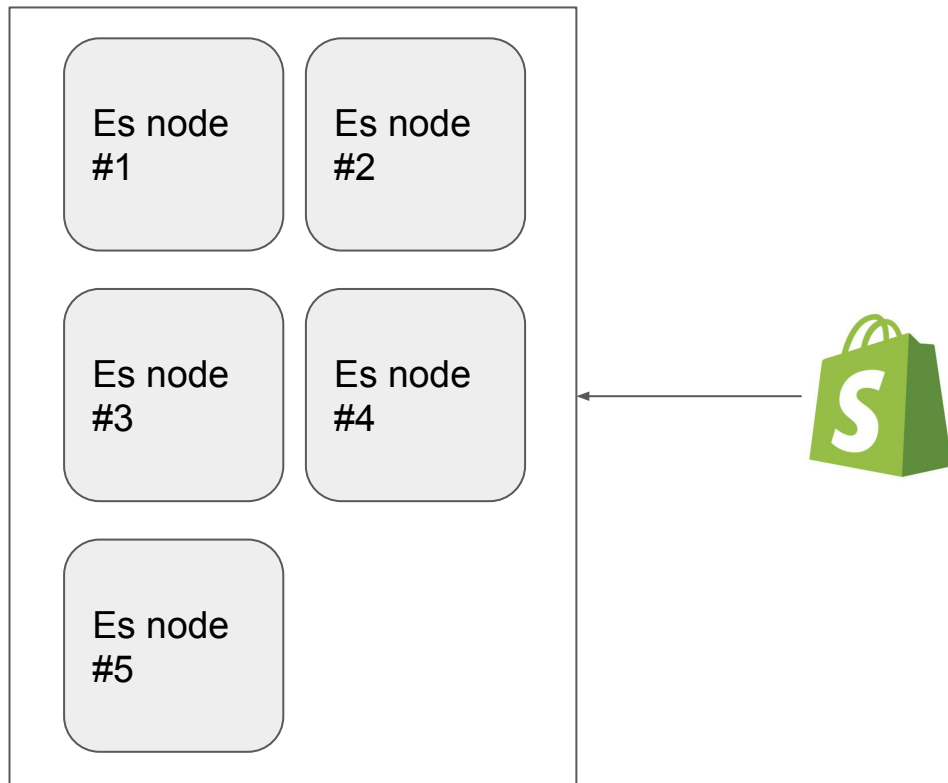
```
apiVersion: stable.shopify.io/v1
kind: Elasticsearch
metadata:
  name: app
spec:
  elastic-search-version: 6
  zones:
    - us-east1-b
    - us-east1-c
    - us-east1-d
  node-specs:
    - replicas: 3
      cpu-limit: "1"
      mem-limit: 2G
      data-volume-size: 10Gi
  snapshot:
    bucket-name: shopify-app-prod-es-snapshots
```

Good practices are ~~not~~ always in place



Running Es with k8s

Elasticsearch cluster



Elasticsearch Kubernetes components

- Statefulset
 - Elasticsearch nodes
- Deployment
 - Snapshot pod with crontab and cli
- Configmap
 - Shared Elasticsearch and logging configuration
- Service
 - Client Endpoint: elasticsearch:9200
 - Discovery
 - Data

Multi-zone data resiliency

- Regions and Zones
- One Statefulset per zone
 - Persistent volume claim
- Allocation awareness setting

North America



northamerica-northeast1

'northamerica-northeast1-a'

'northamerica-northeast1-b'

'northamerica-northeast1-c'

us-west1

'us-west1-a'

'us-west1-b'

'us-west1-c'

us-central1

'us-central1-a'

'us-central1-b'

'us-central1-c'

'us-central1-f'

us-east1

'us-east1-b'

'us-east1-c'

'us-east1-d'

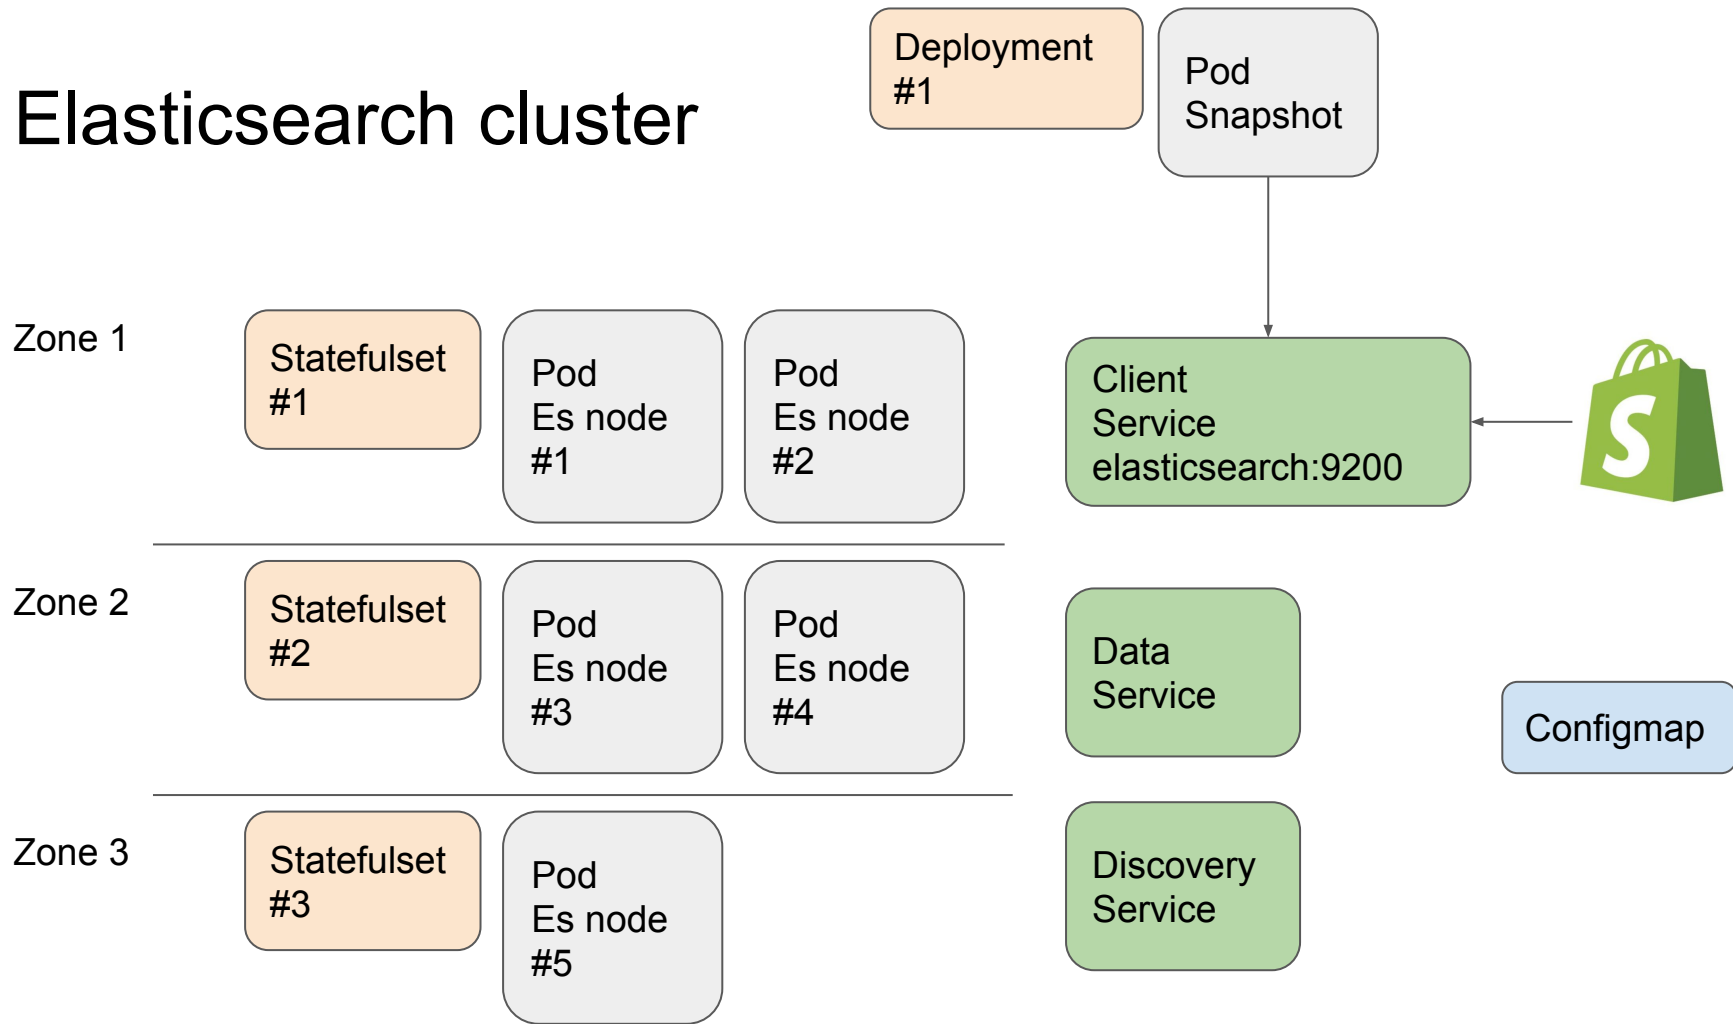
us-east4

'us-east4-a'

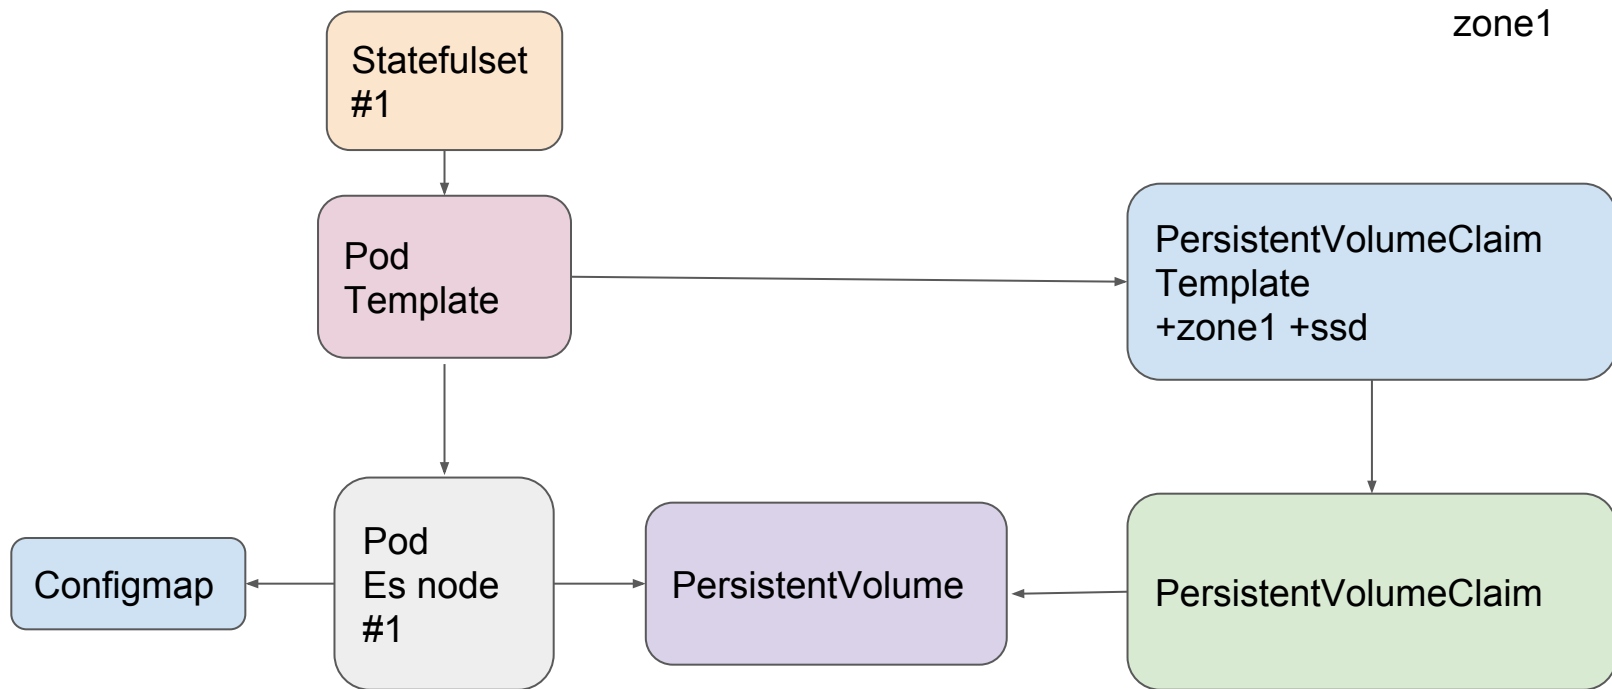
'us-east4-b'

'us-east4-c'

Elasticsearch cluster



Elasticsearch node



```
$ kubectl get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/app--elasticsearch-snapshot-cronjob	1	1	1	3m	

NAME	DESIRED	CURRENT	AGE
statefulsets/app--elasticsearch-data-us-east1-b	1	1	3m
statefulsets/app--elasticsearch-data-us-east1-c	1	1	3m
statefulsets/app--elasticsearch-data-us-east1-d	1	1	3m

NAME	READY	STATUS	RESTARTS	AGE
po/app--elasticsearch-data-us-east1-b-0	1/1	Running	0	3m
po/app--elasticsearch-data-us-east1-c-0	1/1	Running	0	3m
po/app--elasticsearch-data-us-east1-d-0	1/1	Running	0	3m
po/app--elasticsearch-snapshot-cronjob-7d699dd658-f7p8w	1/1	Running	0	3m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/app--elasticsearch	ClusterIP	10.187.243.214	<none>	9200/TCP	3m
svc/app--elasticsearch-data	ClusterIP	10.187.255.181	<none>	9300/TCP	3m
svc/app--elasticsearch-discovery	ClusterIP	10.187.242.125	<none>	9300/TCP	3m

Elasticsearch operations are manual



Automate operations

Rolling update

- When Statefulset or Configmap change:
 - For each node from the previous configuration:
 - Wait for ready/green cluster
 - Stop allocation
 - Restart pod
 - Start allocation
- Downtime on master restart, no native way around
 - Elastic notification of maintenance email > This work will have minimum impact on your usage of the cluster. This change will involve a grow and shrink action on the cluster which you may notice and it may cause a reelection of the master node.

Scaling

- When configured replicas not the same as current n of nodes:
 - While replicas \neq n of nodes
 - Wait for the cluster to be ready and green
 - Increment or decrement one by one
- Minimum master nodes
 - Split brain on scale up
 - Master lost on scale down

Elasticsearch operations are ~~manual~~
automated



Snapshots for everyone

Snapshots

- Bucket
- Crontab
- CLI
- Purge

The big picture

Search operator: Datadog and Bugsnag

Statefulset rolling restart (one by zone)

Show

1h

The Past Hour



Elasticsearch: Datadog, Splunk and Kibana

Used cpu in % (from es API)

Show

1h

The Past Hour



100

90

80

70

60

50

40

30

20

10

0

12:00

Mon 26

12:00

Tue 27


12:00

Wed 28

12:00

March

12:00



2018

One year later

- Managed Elasticsearch clusters: XX+
- Kubernetes pods: XXX+



- Clusters are up to date
- Many automated operations
- Good practices in place

Demo

- Create
- Update
- Scale
- Delete

Conclusion

2017

Scaled search infra with k8s

2018

Search infra as a product (internally)

Thanks!

<https://github.com/maxboisvert/waq2018>

